

A Short Note on Surveillance Testing

FD, CG, VK & NQ on behalf of LMH Governing Body

14/09/2020

Controlling the spread of coronavirus infection within the university requires a reliable means of detecting clusters of infection, and guiding future policy changes, e.g. restricting in-person instruction, demands accurate information about the prevalence of infection. This document outlines some simple strategies to achieve both goals. We focus on detecting the presence of infected *household groups* of students in college or university-provided residential accommodation. We restrict attention to students rather than the university community as a whole because, given limited testing and logistical capacity, it makes sense to focus on the sub-population in which transmission is most likely to occur. We treat household groups, rather than individual students, as our unit of analysis for two reasons: first because this allows us to employ a simple and efficient pooled testing strategy, and second because it is transmission *between* rather than within households that is of more serious concern. The problem of estimating infection prevalence among households and detecting clusters of infection is fundamentally different from that of determining whether a particular individual has the coronavirus. In particular, the high test specificity and sensitivity along with rapid turnaround that are crucial for individual medical diagnosis are not absolutely essential for surveillance testing. Our calculations suggest that relatively modest resources may be sufficient to obtain valuable information, separate from and parallel to the “Early Alert Testing” service. The ideas we discuss are neither novel nor original, and we do not propose a detailed plan for surveillance testing. Our goals are primarily to explain randomized and pooled testing strategies to a non-specialist audience, and to calculate the approximate scale of testing required. We begin by abstracting away the problem of false negatives and false positives. The final section discusses how these affect our analysis.

The Bottom Line

We view a policy of universal pooled testing, like that planned by the University of Cambridge, as the first-best policy for identifying clusters of infection and understanding changes in coronavirus prevalence among our students. Sufficient laboratory capacity for roughly 1500 tests per week would allow fortnightly universal testing of students in residential accommodation, pooling tests within household group. Because capacity cannot be diverted from the Early Alert Testing service, the university would need to obtain around 12,000 tests from an alternative source over the course of MT. These need not have a rapid turnaround

time: a delay of several days would still allow us to assess the spread of infection. If insufficient laboratory capacity is available for a policy of universal testing, randomized testing on the scale of around 250 tests per week (2000 over the course of the term) could provide valuable information about changes in prevalence and, combined with an adaptive sampling scheme, help us to identify clusters of infection as they emerge.

Pooled Testing By Household

Pooled testing is a procedure that can be helpful in situations where laboratory capacity, rather than the ability to collect test samples, is the limiting factor in expanding testing. In this procedure, a single laboratory test is carried out on the *combined* samples of a group of individuals. For a sufficiently accurate test, a negative result for the pooled sample indicates that no one in the group is infected while a positive result indicates that at least one person in the group is infected. Traditionally, a positive group test result would be followed by further testing to determine which individual or individuals in the group is infected. If, however, we are only interested in detecting infected *groups*, then no further testing is required. This makes pooled testing an extremely efficient procedure when applied to household groups: only one test is required to determine whether any members of a given household are infected with coronavirus. Given that household members would be required to isolate if any of them tested positive, for the purposes of reducing the spread of infection it is irrelevant precisely which of them is infected.

There are approximately 24,000 students at the University of Oxford divided roughly evenly between undergraduate and graduate students.¹ As a rough approximation, suppose that 100% of undergraduates and roughly 50% of graduate students are housed in college or university-provided accommodation. With a household size of 6, this makes 3000 households in total. Accordingly, sufficient laboratory capacity for 1500 tests per week would allow fortnightly universal surveillance testing, pooling tests by household. Pooled testing reduces the number of laboratory tests required but does not affect the number of samples that need to be collected. To reduce the burden of collecting samples from 200 students daily, swabs could be self-administered under the guidance of a healthcare professional.²

We view universal pooled testing of households as the first-best policy for understanding and controlling the spread of infection. Should there be insufficient testing capacity to implement such a policy, however, there are alternative strategies that can provide valuable information with fewer resources. The remainder of this document explains some simple and effective possibilities.

¹For precise student numbers, see <https://www.ox.ac.uk/about/facts-and-figures/student-numbers>.

²This is the procedure used by the ONS in their [Coronavirus Infection Survey pilot](#). A recent study suggests that this procedure is no less accurate than having healthcare workers conduct the swabs directly: <https://www.medrxiv.org/content/10.1101/2020.04.11.20062372v1>.

Random Sampling of Households

Suppose that C of the roughly 3000 households in the university contain at least one member who is infected with coronavirus. With a perfect test, we could learn C exactly with 3000 tests, pooled by household. But what if we only have the capacity to carry out n tests, where n is much smaller than 3000? In this case we cannot learn C exactly, but by employing *randomized testing* we can produce an estimate that is sufficiently accurate to guide our mitigation policies and help us to identify clusters of infection. The value of randomized testing is widely acknowledged. It underlies both the [Coronavirus Infection Survey pilot](#) conducted by the ONS, and the [REACT-1](#) survey conducted by the Department of Health and Social Care. Both of these surveys are extremely valuable, but as they aim to be nationally representative, neither employs a sufficiently large sample size to track local trends. As such we propose that randomized surveillance testing be carried out within the university.

The most basic form of randomized testing relies on a *simple random sample*. To use this procedure, we assign each household an id number, place all 3000 id numbers into a hat, mix them up thoroughly, and then draw out n at random. Each of these n households is sent for a pooled coronavirus test. We then use the number of households that test positive for coronavirus to calculate an estimate E of the number households in the population that are infected. For example, if one household in a sample of 300 tests positive, then we would estimate that ten households in the university have coronavirus. Repeating this process weekly would allow us to estimate how the prevalence of coronavirus changes over time.

Because households are chosen at random, the estimated number of infected households E may not equal the true number of infected households C . Purely by chance, we might draw a sample of households in which no one is infected. In this case E would equal zero: an underestimate of C . We could also be unlucky in the opposite direction and, purely by chance, draw a sample that contains everyone at Oxford who is infected. In this case E would be an overestimate of C . Both of these possibilities, however, are remote: E is highly likely to be close to C . With an appropriate sample size, an estimate based on random testing is reliable enough to tell us with high confidence whether C is, say, greater than 50 or fewer than 20. It can also be used to determine whether prevalence is increasing and, if so, how quickly.

Testing Symptomatic Individuals Does Not Suffice

Relative to the Early Alert Testing system, which tests only symptomatic individuals, a crucial advantage of randomized testing, or universal testing, is that it will also detect pre-symptomatic and asymptomatic cases. This is important because, as explained in the [SAGE report of 3rd September](#), “asymptomatic transmission is a key risk in university settings.” Indeed, a recent study cited in the SAGE report finds that only 18% of those aged 0–19 and 22% of those aged 20–39 who are infected with coronavirus show symptoms of the disease.³ Even if these figures are substantial overestimates, it is likely that a considerable fraction of infections among our students will go undetected by the Early Alert Testing system. While

³<https://arxiv.org/abs/2006.08471>

students themselves are at low risk of serious complications from coronavirus, members of staff are at substantially higher risk. By the time the Early Alert system registers an increase in cases among members of staff, it may already be too late for mitigation policies, such as curtailing in-person instruction, to be effective. Even if their turn-around times were somewhat slower than those for tests conducted through the Early Alert system, randomized pooled tests could still provide more timely information about the spread of infection, given the likely extent of asymptomatic transmission among our students.

How many samples are required?

The key question in designing a randomized testing protocol is how large a sample size to use. To answer this question, we first need a rough idea of the prevalence of coronavirus among households at the beginning of term. Based on recent results from [REACT-1](#), approximately 0.25% of individuals in the 18-24 age group had the coronavirus at the start of September. Taking this as a lower bound for prevalence in early October, suppose that at least 45 of the roughly 18,000 students in residential accommodation has coronavirus at the start of MT. Given that our interest is in households rather than individuals, the question remains: if 45 students have coronavirus, then how many *households* are infected?

The answer depends on what we are willing to assume about households. At one extreme, as many as 45 households could be infected: this would require that exactly *one* person in each affected household has the coronavirus. At the other extreme, it is possible that as few as 8 households could be infected. This would require 7 households in which every member was infected, and 1 household with 3 infected members. During term, intra-household transmission will likely be substantial.⁴ At the beginning of term, however, it is unlikely that cases will cluster strongly within households, as most groups will not have met in person over the summer. If initial coronavirus cases can be viewed as independently and uniformly assigned to households, then 45 infected students would translate into approximately 44.7 infected households on average.⁵ To allow for some initial clustering of cases, lower this by a third and suppose that there are initially 35 infected households, yielding a household prevalence of approximately 1.16% at the start of term.

Given a rough idea of household prevalence at the start of MT, there are several simple ways to approximate the sample size required for randomized pooled testing to provide useful information. The first and simplest is to ask how many households we would need to test, on average, before obtaining our first positive result. If C out of the 3000 households are infected at the start of term, then the answer is approximately $3000/C$.⁶ Table 1 provides exact results over a range of values for C . This calculation provides a rough-and-ready method for showing that a proposed sample size is likely *too small* to be useful. For example, a sample

⁴A secondary infection rate of 38% within households would be consistent with a recent study of patients in New York State: <https://doi.org/10.1093/cid/ciaa549>.

⁵Specifically, if N is the number of households, k is household size, and S is the number of infected students, the expected number of infected households is $N \left[1 - \binom{kN-S}{k} / \binom{kN}{k} \right]$.

⁶The exact value is $(N+1)/(C+1)$ for a population of size N containing C positives.

size of 50 households would be too small to provide useful information when C is close to 35 because we would, on average, require approximately 83 samples to obtain our first positive result if this were the true prevalence.

Table 1: Expected number of samples before the first positive result under random sampling without replacement from a population of size 3000.

Infected Households (C)	Prevalence (%)	Expected #Samples
25	0.83	115
30	1.00	97
35	1.17	83
40	1.33	73
45	1.50	65
50	1.67	59

A more refined way to determine an appropriate sample size is by calculating *how close* the estimated number of infected households, E , will likely be to the true number of infected households C for different particular values of n .⁷ Suppose, as before, that there are initially 35 infected households. With a sample size of $n = 250$ the estimated number of infected households will fall in the range $[12, 60]$, inclusive, with approximately 90% probability.⁸

Update power calculation to Fisher’s exact test! The population size is too small for a normal approximation to be reasonable here.

A third way of deciding how many samples are needed is by calculating the chance that we will be able to reliably distinguish between two different levels of coronavirus prevalence given a particular sample size. Suppose, for example, that we initially believed there to be 50 cases in the university, corresponding to the current estimated prevalence among the university-aged population from [REACT-1](#). If cases increased to 100, we would have approximately a 68% chance of detecting an increase, with high confidence, based on a sample 1400 students.⁹ If they increased to 200, we would have a 99% chance of detecting an increase. Roughly speaking, if cases were doubling weekly, then a sample of 200 students per day would be likely to detect an increase after one week, and nearly certain to do so after two weeks.

Broadly speaking, our calculations suggests that a simple random sample of approximately 250 household groups per week to estimate coronavirus prevalence at the household level and to determine whether, and how quickly, it may be increasing. This would amount to around

⁷This interval is calculated from the quantiles of a Hypergeometric(N, C, n) distribution where N is the population size, C is the number of positives in the population, and n is the sample size.

⁸This interval is calculated from the quantiles of a Hypergeometric(N, C, n) distribution where N is the population size, C is the number of positives in the population, and n is the sample size.

⁹These values correspond to a power calculation for an $\alpha = 0.1$ test of $H_0: p \leq p_0$ versus $H_1: p > p_0$ using the score test statistic $T_n = (\hat{p} - p_0)/SE(p_0)$ where $SE(p_0) = \sqrt{\kappa p_0(1 - p_0)/n}$ and $\kappa = (N - n)/(N - 1)$ is a finite-population correction. In finite sample simulations, a standard normal critical value yields a slightly over-sized test, $\alpha \approx 0.13$ rather than 0.1, resulting in slightly higher power than this calculation would suggest.

??? as many tests over the entire term as would be required for a *single* round of universal testing.

Finding Clusters with Adaptive Sampling

Thus far assumed we want to estimate prevalence and considered simplest possible sampling procedure: simple random sample. May also want to try to find and isolate clusters of infection. In fact a more sophisticated sampling design can achieve both goals simultaneously: adaptive cluster sampling.¹⁰ Define “neighborhoods” of households between which transmission is particularly likely, e.g. residence halls. If one household tests positive, immediately test the other households in its neighborhood. Deviation from simple random sample is more complicated to analyze mathematically but blah blah blah very good blah blah.¹¹

What about false positives and negatives?

Thus far we have abstracted away the problem of false positives and false negatives. As our primary goal is to estimate prevalence and identify clusters of infection rather than to diagnose particular individuals, this is not a major concern. First, if false positive and false negative rates are constant, we can still gain useful information by examining *changes* in estimated prevalence. Second, recent research has proposed estimates of these rates that can be used to correct estimates of population prevalence.

¹⁰Thompson, S.K. (1990), “Adaptive Cluster Sampling”, *Journal of the American Statistical Association*, 85(412), pp. 1050-1059.

¹¹Thompson, S.K. (1990), “Adaptive Cluster Sampling”, *Journal of the American Statistical Association*, 85(412), pp. 1050-1059.